



400 x 400 x 3 convolve this with 3 x 3 x 3 x 32

What size of image do we really need?



The 8 frames we'll send — note how little changes between them

t=8.2s

t=16.4s

t=24.6s

t=32.8s

t=41.0s

t=49.2s

t=57.4s



Same frame, fewer pixels — still legible, far cheaper to send


512px
15.1 KB

256px
5.8 KB

128px
2.6 KB

64px
1.4 KB





```
[{'type': 'text',
  'text': "TRANSCRIPT (timestamped):\n[0.0-1.0s] Lieutenant Drebin, police squad.\n[1.0-2.0s] Oh,
{'type': 'text', 'text': '[frame @ 0.0s]'},
{'type': 'image_url',
  'image_url': {'url': 'data:image/jpeg;base64,/9j/4AAQSkZJRgABAQAAQABAAD/2wBDAAYEBQYFBAYGBQYHBwY
  'detail': 'low'}},
{'type': 'text', 'text': '[frame @ 16.4s]'},
{'type': 'image_url',
  'image_url': {'url': 'data:image/jpeg;base64,/9j/4AAQSkZJRgABAQAAQABAAD/2wBDAAYEBQYFBAYGBQYHBwY
  'detail': 'low'}},
{'type': 'text', 'text': '[frame @ 32.8s]'},
{'type': 'image_url',
  'image_url': {'url': 'data:image/jpeg;base64,/9j/4AAQSkZJRgABAQAAQABAAD/2wBDAAYEBQYFBAYGBQYHBwY
  'detail': 'low'}},
{'type': 'text', 'text': '[frame @ 49.2s]'},
{'type': 'image_url',
  'image_url': {'url': 'data:image/jpeg;base64,/9j/4AAQSkZJRgABAQAAQABAAD/2wBDAAYEBQYFBAYGBQYHBwY
  'detail': 'low'}},
{'type': 'text', 'text': 'What movie it is coming from?'}]
```

Entire transcript:

↓ Jm

↓ Jm

↓ Jm

↓ Jm

N



↓
Jm →

Que
↓
Jm

CRITIC

REFLECT

TUBELET

Chroma DB (online): one kind of embedding per chunk (simplistic RAG Approaches)

PineCONE DB: very scalable and available.

QDRANT(super vector DB): Multiple kind of embedding per chunk, with more advanced filtering conditions per chunk